Query Optimization

Reduction Factor

1. Col = value:

$$RF = 1/NKeys(Col)$$

2. Col > value:

$$RF = (High(Col) - value)/(High(Col) - Low(Col))$$

3. Col < value:

$$RF = (val - Low(Col))/(High(Col) - Low(Col))$$

4. ColA = ColB (for joins): ??

$$RF = 1/(\max(NKeys(ColA), NKeys(ColB)))$$

5. No information about Nkeys or intervals

$$RF = 1/10$$

Result Size Estimation

1. Single table selection:

$$ResultSize = NTuples(R) \times \prod_{i=1,..,n} RF_i$$

2. Joins (over k tables):

$$ResultSize = \prod_{j=1,...,k} NTuples(R_j) \prod_{i=1,..,n} RF_i$$

if no selections (no predicates), $RF = 1$

Single-relation plans

1. Sequenctial (heap) scan:

$$Cost = NPages(R)$$

2. Index selection over a primary key (single tuple):

$$Cost(B + Tree) = Height(I) + 1$$

$$Cost(HashIndex) = ProbeCost(I) + 1, \quad ProbeCost(I) \approx 1.2$$

3. Clustered index matching one or more predicates:

$$Cost(B + Tree) = (NPages(I) + NPages(R)) \times \prod_{i=1,...n} RF_i$$

$$Cost(HashIndex) = 2.2 \times NPages(R) \times \prod_{i=1,...n} RF_i$$

4. Non-clustered index ...

$$Cost(B + Tree) = (NPages(I) + NTuples(R)) \times \prod_{i=1,...n} RF_i$$

$$Cost(HashIndex) = 2.2 \times NTuples(R) \times \prod_{i=1,...n} RF_i$$

Multi-relation Plans

Step-by-step:

1. Select order of relations: $S \times R \times B,\ S \times B \times R, \ldots \quad \Rightarrow N!$
2. For each join, select join algorithm: Hash join, Sort-merge join...
3. For each input relation, select access method: Heap scan, various index alternatives..
4. Calculate ResultSize and Costs for each step, then compute the total cost

Onluy left-deep joint trees are considered: Intermediate results are not written to temporary files

Example: